ECE408/CS483/CSE408 Fall 2022

Applied Parallel Programming

Lecture 5: Locality and Tiled Matrix Multiplication

Course Reminders

- We are grading Lab 1 now
- Lab 2 is out; it is due this Friday
- Lowest lab grade will be dropped from the final grade
 - Thus, no late submissions are allowed for labs

Objective

- To learn to evaluate the performance implications of global memory accesses
- To prepare for MP3: tiled matrix multiplication
- To learn to assess the benefit of tiling

Kernel Invocation (Host-side Code)

```
// Setup the execution configuration
// BLOCK WIDTH is a #define constant
dim3 dimGrid(ceil((1.0*Width)/BLOCK WIDTH),
           ceil((1.0*Width)/BLOCK WIDTH), 1);
dim3 dimBlock(BLOCK WIDTH, BLOCK WIDTH, 1);
// Launch the device computation threads!
MatrixMulKernel<<<dimGrid, dimBlock>>>(Md, Nd, Pd, Width);
```

A Simple Matrix Multiplication Kernel

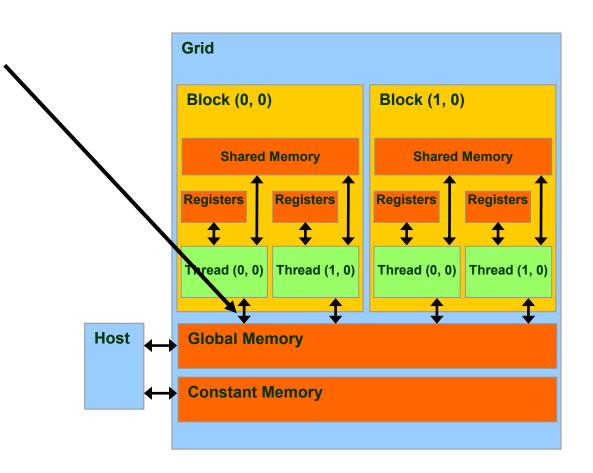
```
global
void MatrixMulKernel(float *d M, float *d N, float *d P, int Width)
   // Calculate the row index of the d P element and d M
   int Row = blockIdx.y*blockDim.y+threadIdx.y;
   // Calculate the column idenx of d P and d N
   int Col = blockIdx.x*blockDim.x+threadIdx.x;
   if ((Row < Width) && (Col < Width)) {
      float Pvalue = 0;
      // each thread computes one element of the block sub-matrix
      for (int k = 0; k < Width; ++k)
          Pvalue += d M[Row*Width+k] * d_N[k*Width+Col];
      d P[Row*Width+Col] = Pvalue;
```

Review: 4B of Data per FLOP

- Each threads access global memory
 - -for elements of M and N:
 - -4B each, or 8B per pair.
 - –(And once TOTAL to P per thread—ignore it.)
- With each pair of elements,
 - -a thread does a single multiply-add,
 - -2 FLOP—floating-point operations.
- So for every FLOP,
 - -a thread needs 4B from memory:
 - **-4B / FLOP**

How about performance on a device with 150 GB/s memory bandwidth?

- All threads access global memory for their input matrix elements
 - Two memory accesses (8 bytes) per floating point multiply-add (2 fp ops)
 - 4B/s of memory bandwidth/FLOPS
 - 150 GB/s limits the code at 37.5 GFLOPS
- The actual code runs at about 25 GFLOPS
- Need to drastically cut down memory accesses to get closer to the peak of more than 1,000 GFLOPS



A Common Programming Strategy

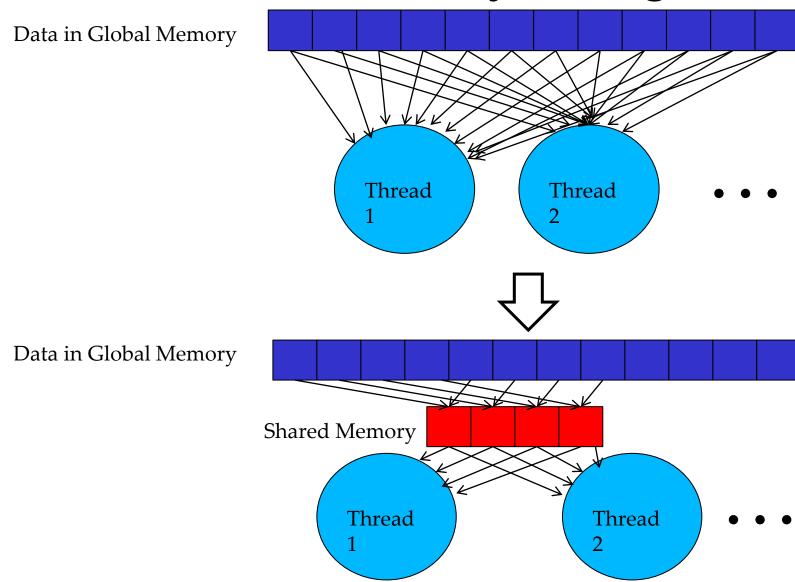
- Global memory is implemented with DRAM slow
- To avoid Global Memory bottleneck, tile the input data to take advantage of Shared Memory:
 - Partition data into subsets (tiles) that fit into the (smaller but faster) shared memory
 - Handle each data subset with one thread block by:
 - Loading the subset from global memory to shared memory, using multiple threads to exploit memory-level parallelism
 - Performing the computation on the subset from shared memory; each thread can efficiently access any data element
 - Copying results from shared memory to global memory
 - Tiles are also called blocks in the literature

A Common Programming Strategy

- In a GPU, only threads in a block can use shared memory.
- Thus, each block operates on separate tiles:
 - Read tile(s) into shared memory using multiple threads to exploit memory-level parallelism.
 - Compute based on shared memory tiles.
 - Repeat.
 - Write results back to global memory.

Declaring Shared Memory Arrays

Shared Memory Tiling Basic Idea



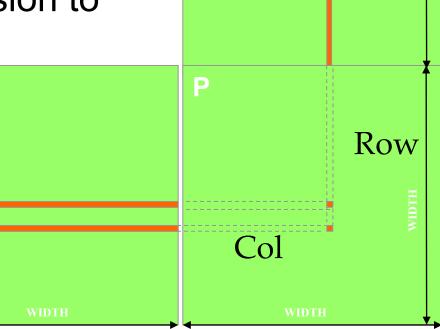
Outline of Technique

- Identify a tile of global data that are accessed by multiple threads
- Load the tile from global memory into on-chip memory
- Have the multiple threads to access their data from the on-chip memory
- Move on to the next block/tile

Use Shared Memory for data that will be reused

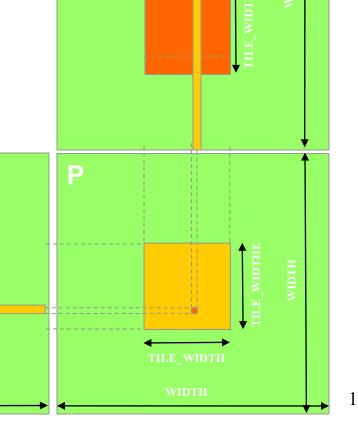
Observe that each input element of M and N is used WIDTH times

 Load each element into Shared Memory and have several threads use the local version to reduce the memory bandwidth



Tiled Multiply

- Break up the execution of the kernel into phases so that the data accesses in each phase are focused on one tile of M and N
- For each tile:
 - Phase 1: Load tiles of M & N into share memory
 - Phase 2: Calculate partial dot product for tile of P



012 TILE WIDTH-1

© David Kirk/NVIDIA and Wen-mei W. Hwu, 2007-2018

ECE408/CS483/ University of Illinois at Urbana-Champaign

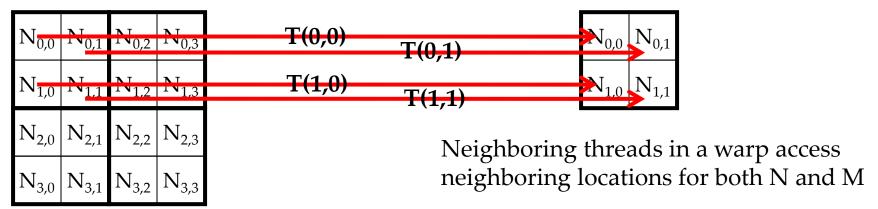
Loading a Tile

- All threads in a block participate
 - Each thread loads
 - one M element and
 - one N element
 - in basic tiling code.

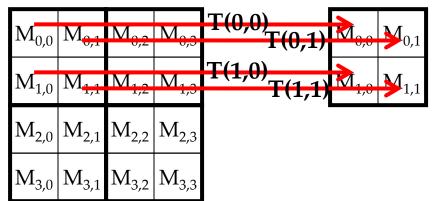
 Assign the loaded element to each thread such that the accesses within each warp is coalesced (more later).

Loading Tiles for Block (0,0)

Shared Memory



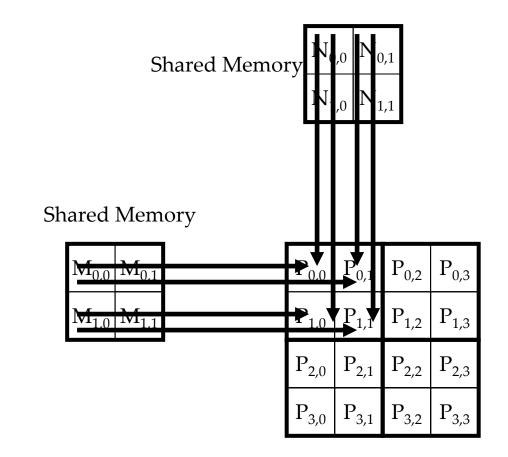
Shared Memory



P _{0,0}	P _{0,1}	P _{0,2}	P _{0,3}
P _{1,0}	P _{1,1}	P _{1,2}	P _{1,3}
P _{2,0}	P _{2,1}	P _{2,2}	P _{2,3}
P _{3,0}	P _{3,1}	P _{3,2}	P _{3,3}

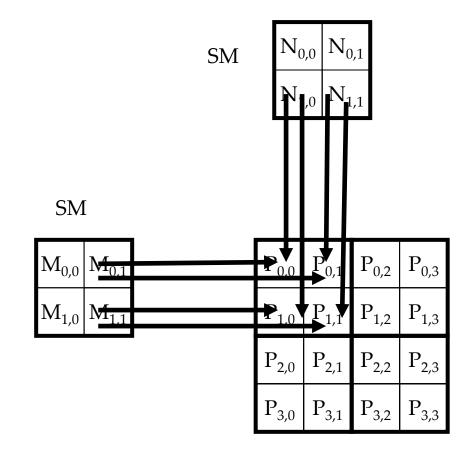
N _{0,0}	N _{0,1}	N _{0,2}	N _{0,3}
N _{1,0}	N _{1,1}	N _{1,2}	N _{1,3}
N _{2,0}	N _{2,1}	N _{2,2}	N _{2,3}
N _{3,0}		N _{3,2}	N _{3,3}

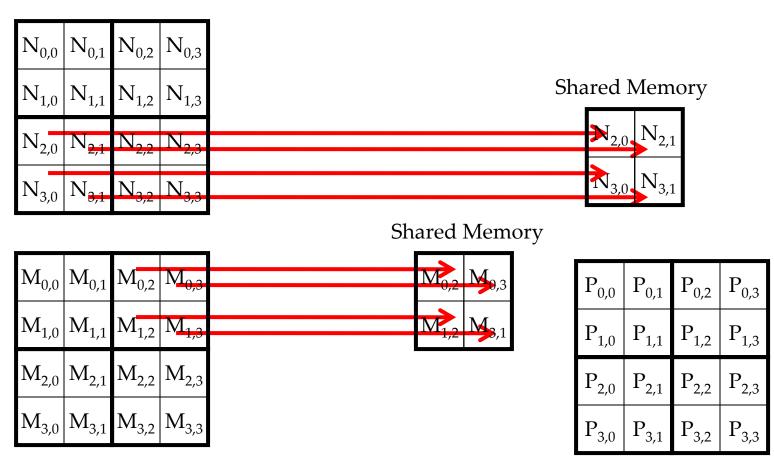
$M_{0,0}$	M _{0,1}	$M_{0,2}$	M _{0,3}
$M_{1,0}$	$M_{1,1}$		
$M_{2,0}$	M _{2,1}	M _{2,2}	M _{2,3}
$M_{3,0}$	M _{3,1}	M _{3,2}	M _{3,3}



N _{0,0}	N _{0,1}	N _{0,2}	N _{0,3}
N _{1,0}	N _{1,1}	N _{1,2}	N _{1,3}
N _{2,0}	N _{2,1}	N _{2,2}	N _{2,3}
N _{3,0}	N _{3,1}	N _{3,2}	N _{3,3}

$M_{0,0}$	$M_{0,1}$	$M_{0,2}$	$M_{0,3}$
M _{1,0}	M _{1,1}	M _{1,2}	M _{1,3}
$M_{2,0}$	M _{2,1}	$M_{2,2}$	$M_{2,3}$
$M_{3,0}$	$M_{3,1}$	$M_{3,2}$	$M_{3,3}$

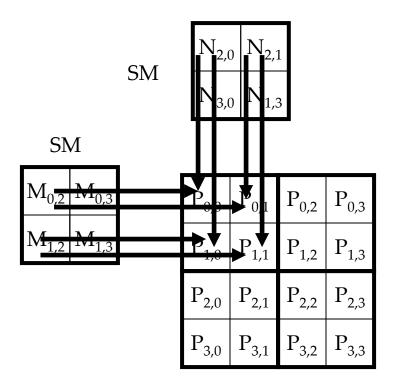




¹⁹

N _{0,0}	N _{0,1}	N _{0,2}	N _{0,3}
N _{1,0}	N _{1,1}	N _{1,2}	N _{1,3}
N _{2,0}	N _{2,1}	N _{2,2}	N _{2,3}
N _{3,0}	N _{3,1}	N _{3,2}	N _{3,3}

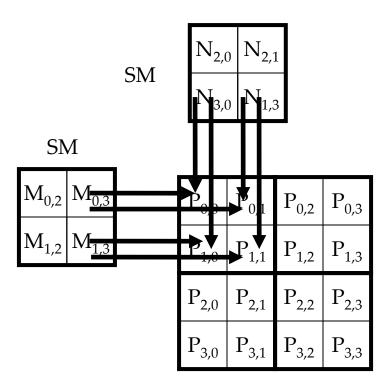
$M_{0,0}$	M _{0,1}	$M_{0,2}$	$M_{0,3}$
$M_{1,0}$	M _{1,1}	M _{1,2}	M _{1,3}
$M_{2,0}$	M _{2,1}	M _{2,2}	$M_{2,3}$
M_{30}	$M_{3.1}$	$M_{3,2}$	$M_{3.3}$



20

N _{0,0}	N _{0,1}	N _{0,2}	N _{0,3}
N _{1,0}	N _{1,1}	N _{1,2}	N _{1,3}
N _{2,0}	N _{2,1}	N _{2,2}	N _{2,3}
N _{3,0}	N _{3,1}	N _{3,2}	N _{3,3}

$M_{0,0}$	M _{0,1}	$M_{0,2}$	$M_{0,3}$
$M_{1,0}$	M _{1,1}	$M_{1,2}$	M _{1,3}
$M_{2,0}$	M _{2,1}	M _{2,2}	M _{2,3}
$M_{2,0}$	$M_{2,1}$	$M_{3,2}$	$M_{3,3}$



Phase 1: Loading a Tile

- All threads in a block participate
 - Each thread loads one M element and one N element in basic tiling code
- Assign the loaded element to each thread such that the accesses within each warp is coalesced (more later).

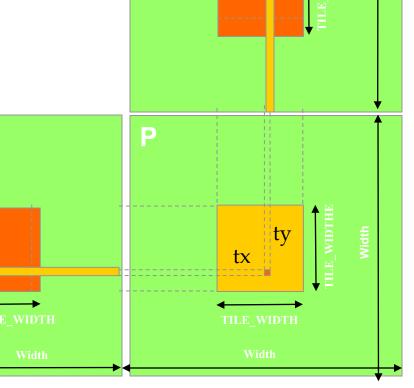
Loading an Input Tile 0

0

TILE WIDTH

2D indexing for Tile 0

M[Row][tx] N[ty][Col]



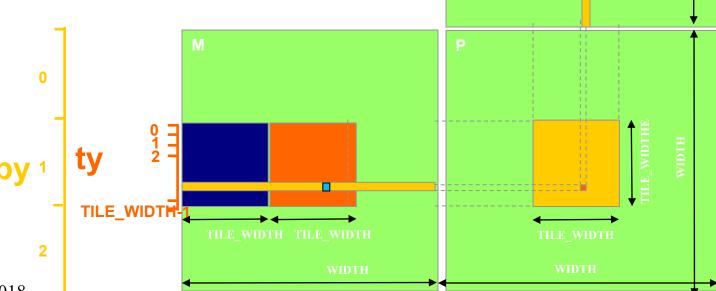
012 TILE WIDTH-1

© David Kirk/NVIDIA and Wen-mei W. Hwu, 2007-2018
ECE408/CS483/ University of Illinois at Urbana-Champaign

23

Loading an Input Tile 1

Accessing tile 1 in 2D indexing:



012 TILE WIDTH-1

24

© David Kirk/NVIDIA and Wen-mei W. Hwu, 2007-2018
ECE408/CS483/ University of Illinois at Urbana-Champaign

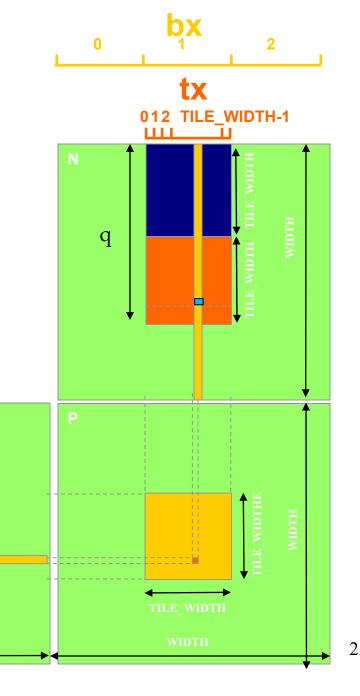
Loading an Input Tile q

However, recall that M and N are dynamically allocated and can only use 1D indexing:

```
M[Row][m*TILE_WIDTH+tx]
M[Row*Width + q*TILE_WIDTH + tx]
```

N[q*TILE_WIDTH+ty][Col]
N[(q*TILE WIDTH+ty) * Width + Col]

TILE WIDTH



© David Kirk/NVIDIA and Wen-mei W. Hwu, 2007-2018

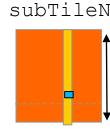
ECE408/CS483/ University of Illinois at Urbana-Champaign

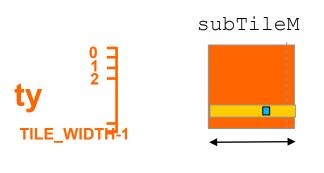
Phase 2: Compute partial product

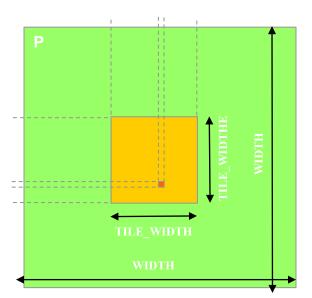
To perform the kth step of the product within the tile:

```
subTileM[ty][k]
subTileN[k][tx]
```









We're Not There Yet!

• But ...

- How can a thread know ...
 - That another thread has finished its part of the tile?
 - Or that another thread has finished using the previous tile?

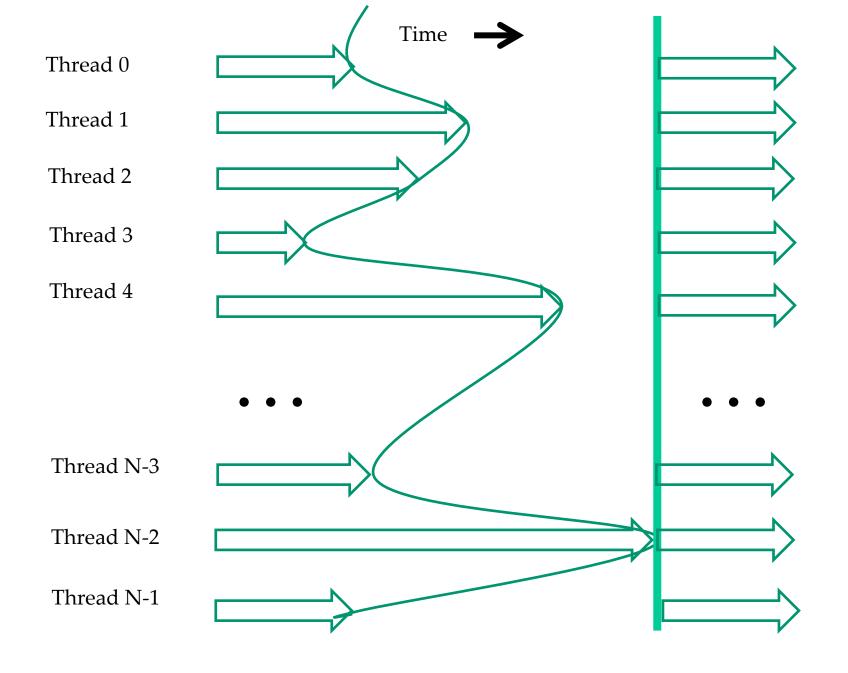
We need to synchronize!

Leveraging Parallel Strategies

- Bulk synchronous execution:
 - threads execute roughly in unison
 - 1. Do some work
 - 2. Wait for others to catch up
 - 3. Repeat
- Much easier programming model
 - Threads only parallel within a section
 - Debug lots of little programs
 - Instead of one large one.
- Dominates high-performance applications

Bulk Synchronous Steps Based on Barriers

- How does it work?
 Use a barrier to wait for thread to 'catch up.'
- A barrier is a synchronization point:
 - each thread calls a function to enter barrier;
 - threads block (sleep) in barrier function until all threads have called;
 - after last thread calls function,
 all threads continue past the barrier.



Barrier Synchronization

- An API function call in CUDA __syncthreads()
- All threads in the same block must reach the ___syncthreads()
 before any can move on

- Can be used to coordinate tiled algorithms
 - To ensure that all elements of a tile are loaded
 - To ensure that certain computation on elements is complete

Tiled Matrix Multiplication Kernel

```
global void MatrixMulKernel(float* M, float* N, float* P, int Width)
    shared float subTileM[TILE WIDTH] [TILE WIDTH];
   shared float subTileN[TILE WIDTH] [TILE WIDTH];
3. int bx = blockIdx.x; int by = blockIdx.y;
4. int tx = threadIdx.x; int ty = threadIdx.y;
   // Identify the row and column of the P element to work on
5. int Row = by * TILE WIDTH + ty;
6. int Col = bx * TILE WIDTH + tx;
   float Pvalue = 0;
   // Loop over the M and N tiles required to compute the P element
   // The code assumes that the Width is a multiple of TILE WIDTH!
8. for (int q = 0; q < Width/TILE WIDTH; ++q) {
      // Collaborative loading of M and N tiles into shared memory
      subTileM[ty][tx] = M[Row*Width + q*TILE WIDTH+tx];
9.
      subTileN[ty][tx] = N[(q*TILE WIDTH+ty)*Width+Col];
10.
11.
      syncthreads();
12.
      for (int k = 0; k < TILE WIDTH; ++k)
13.
          Pvalue += subTileM[ty][k] * subTileN[k][tx];
14.
      syncthreads();
15. }
16. P[Row*Width+Col] = Pvalue;
```

Compare with Basic MM Kernel

```
global void MatrixMulKernel(float* M, float* N, float* P, int Width)
// Calculate the row index of the P element and M
int Row = blockIdx.y * blockDim.y + threadIdx.y;
// Calculate the column index of P and N
int Col = blockIdx.x * blockDim.x + threadIdx.x;
if ((Row < Width) && (Col < Width)) {
  float Pvalue = 0;
   // each thread computes one element of the block sub-matrix
   for (int k = 0; k < Width; ++k)
     Pvalue += M[Row*Width+k] * N[k*Width+Col];
   P[Row*Width+Col] = Pvalue;
```

Use of Large Tiles Shifts Bottleneck

- Recall our example GPU: 1,000 GFLOP/s, 150 GB/s
- 16x16 tiles use each operand for 16 operations
 - reduce global memory accesses by a factor of 16
 - 150GB/s bandwidth supports (150/4)*16 = 600 GFLOPS!
- 32x32 tiles use each operand for 32 operations
 - reduce global memory accesses by a factor of 32
 - -150 GB/s bandwidth supports (150/4)*32 = 1,200 GFLOPS!
 - Memory bandwidth is no longer the bottleneck!

Also Need Parallel Accesses to Memory

- Shared memory size
 - implementation dependent
 - 64kB per SM in Maxwell (48kB max per block)
- Given TILE_WIDTH of 16 (256 threads / block),
 - each thread block uses2*256*4B = 2kB of shared memory,
 - which limits active blocks to 32;
 - max. of 2048 threads per SM,
 - which limits blocks to 8.
 - Thus up to 8*512 = 4,096 pending loads
 (2 per thread, 256 threads per block)

Another Good Choice: 32x32 Tiles

- Given TILE_WIDTH of 32 (1,024 threads / block),
 - each thread block uses2*1024*4B = 8kB of shared memory,
 - which limits active blocks to 8;
 - max. of 2,048 threads per SM,
 - which limits blocks to 2.
 - Thus up to 2*2,048 = **4,096 pending loads** (2 per thread, 1,024 threads per block)

(same memory parallelism exposed)

Current GPU? Use Device Query

Number of devices in the system

```
int dev_count;
cudaGetDeviceCount( &dev_count);
```

Capability of devices

```
cudaDeviceProp dev_prop;
for (i = 0; i < dev_count; i++) {
          cudaGetDeviceProperties( &dev_prop, i);
          // decide if device has sufficient resources and capabilities
}</pre>
```

- cudaDeviceProp is a built-in C structure type
 - dev_prop.dev_prop.maxThreadsPerBlock
 - Dev_prop.sharedMemoryPerBlock

– ...

ANY MORE QUESTIONS? READ CHAPTER 4!