



ECE408/CS483/CSE408 Fall 2022

Applied Parallel Programming

# Lecture 23: Alternatives to CUDA

# Accelerated Computing is no longer a question



GPU vendors include:

Nvidia

AMD

Intel

Samsung

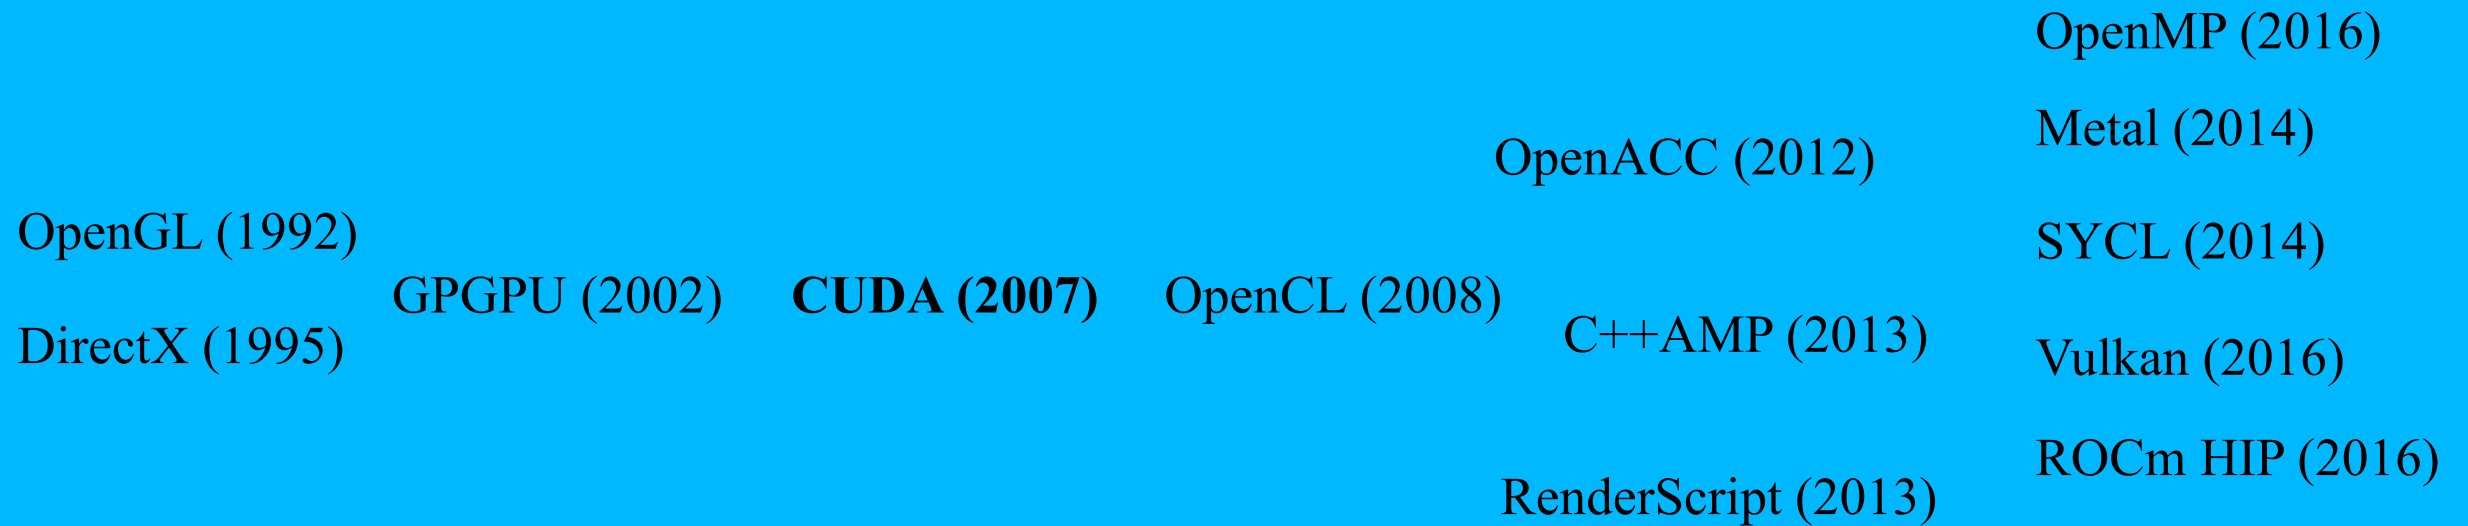
Apple

Qualcomm

ARM

Etc....

# CUDA is just one model for Compute Acceleration



Existing frameworks such as MPI, TBB, OpenCV adapted to provide support.

New frameworks such as Caffe, TensorFlow, R, PyCUDA natively support acceleration.

# OpenCL, HIP, OpenACC, MPI

- OpenCL: An Open Standard Acceleration API
- Heterogeneous-Computing Interface for Portability (HIP)
- OpenACC: A “Low-Code” Acceleration API
- MPI: A Large Scale, Multi-Node Parallel API

# Common Traits for Acceleration APIs

- HARDWARE

- Hierarchy of lightweight cores
- Local scratchpad memories
- Lack of HW coherence
- Slow global atomics
- Threading

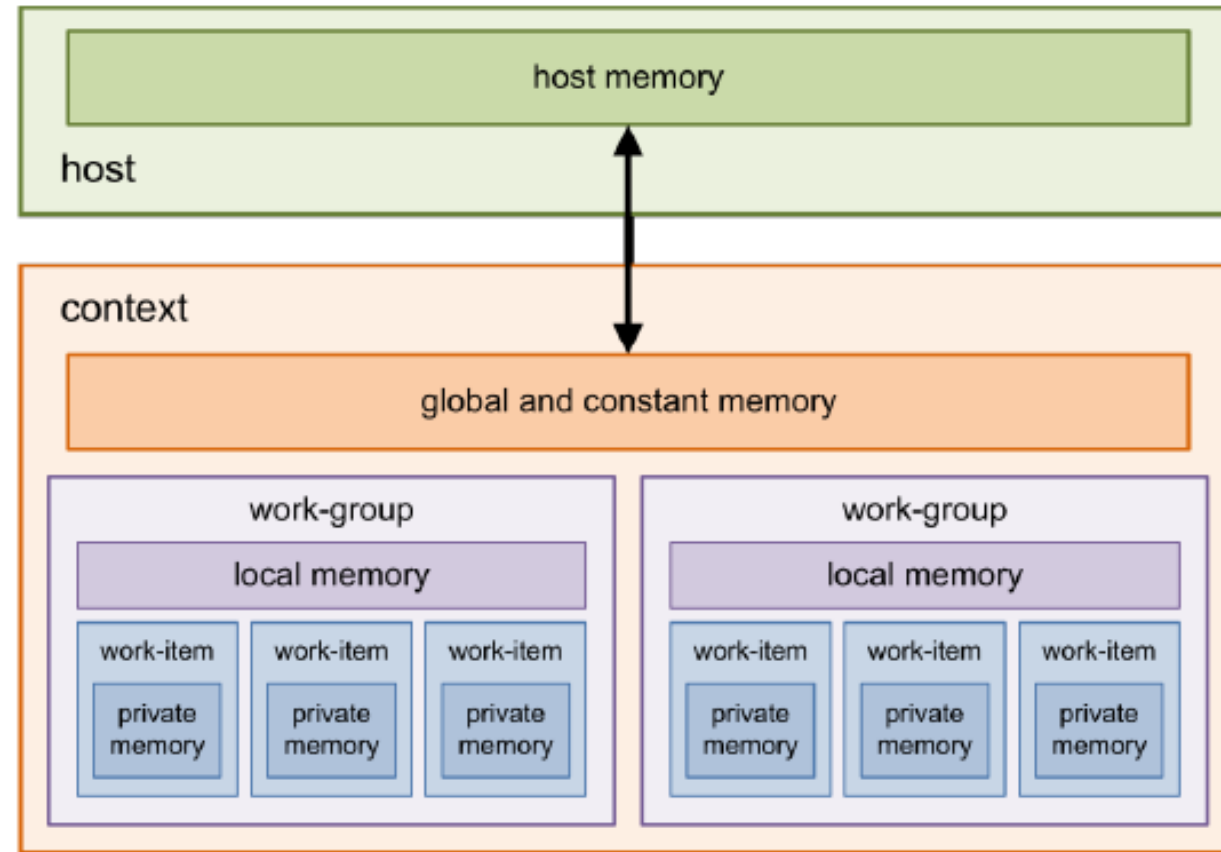
- SOFTWARE

- Kernel oriented acceleration
- Device memory vs. Host memory
- Software managed memory
- Grids, Blocks, Threads
- Bulk Synchronous Parallelism

# OpenCL

- Framework for CPUs, GPUs, DSPs, FPGAs, etc. (not just Nvidia GPUs)
- Initially developed by Apple with support from AMD, IBM, Qualcomm, Intel, and Nvidia. OpenCL 1.0 launched in 2008.
- OpenCL 2.2 launched in May 2017
- Apple announces dropping of OpenCL in 2018

# OpenCL Memory Model



# OpenCL MatMult

- Notice similarity to CUDA
- WorkGroup similar to Block
- WorkItem similar to Thread
- `__local` similar to `__shared`

```
1.// Tiled and coalesced version
2. __kernel void myGEMM2(int M, int N, int K, __global float* A, __global float* B, __global float*
3.) {
4.     // Thread identifiers
5.     const int row = get_local_id(0); // Local row ID (max: TS)
6.     const int col = get_local_id(1); // Local col ID (max: TS)
7.     const int globalRow = TS*get_group_id(0) + row; // Row ID of C (0..M)
8.     const int globalCol = TS*get_group_id(1) + col; // Col ID of C (0..N)
9.
10.    // Local memory to fit a tile of TS*TS elements of A and B
11.    __local float Asub[TS][TS];
12.    __local float Bsub[TS][TS];
13.
14.    // Initialise the accumulation register
15.    float acc = 0.0f;
16.    // Loop over all tiles
17.    const int numTiles = K/TS;
18.    for (int t=0; t<numTiles; t++) {
19.
20.        // Load one tile of A and B into local memory
21.        const int tiledRow = TS*t + row;
22.        const int tiledCol = TS*t + col;
23.        Asub[col][row] = A[tiledCol*M + globalRow];
24.        Bsub[col][row] = B[globalCol*K + tiledRow];
25.
26.        // Synchronise to make sure the tile is loaded
27.        barrier(CLK_LOCAL_MEM_FENCE);
28.
29.        // Perform the computation for a single tile
30.        for (int k=0; k<TS; k++)
31.            acc += Asub[k][row] * Bsub[col][k];
32.
33.        // Synchronise before loading the next tile
34.        barrier(CLK_LOCAL_MEM_FENCE);
35.    }
36.
37.    // Store the final result in C
38.    C[globalCol*M + globalRow] = acc;
39.}
```



# HIP

- Heterogeneous-Computing Interface for Portability (HIP)
  - C++ dialect designed to ease conversion of CUDA applications to portable C++ code.
  - Provides a C-style API and a C++ kernel language.
  - The C++ interface can use templates and classes across the host/kernel boundary.
- HIP code can run on AMD hardware (through the HCC compiler) or NVIDIA hardware (through the NVCC compiler).
- The HIPify tool automates much of the conversion work by performing a source-to-source transformation from CUDA to HIP.

# vectorAdd with HIP

```
__global__ void vecAdd(double *a, double *b, double *c, int n) {
    int id = blockIdx.x*blockDim.x+threadIdx.x;
    if (id < n) c[id] = a[id] + b[id];
}
...
hipMalloc(&d_a, nbytes);
hipMalloc(&d_b, nbytes);
hipMalloc(&d_c, nbytes);

hipMemcpy(d_a, h_a, bytes, hipMemcpyHostToDevice);
hipMemcpy(d_b, h_b, bytes, hipMemcpyHostToDevice);

blockSize = 1024;
gridSize = (int)ceil((float)n/blockSize);

hipLaunchKernelGGL(vecAdd, dim3(gridSize), dim3(blockSize), 0, 0, d_a, d_b, d_c, n);
hipDeviceSynchronize( );

hipMemcpy(h_c, d_c, bytes, hipMemcpyDeviceToHost);
...
```

# OpenACC

The OpenACC Application Programming Interface (API) provides a set of

- compiler directives (pragmas),
- library routines, and
- environment variables

that enable

- FORTRAN, C and C++ programs
- to execute on accelerator devices
- including GPUs and CPUs.

# Pragmas Provide Extra Information

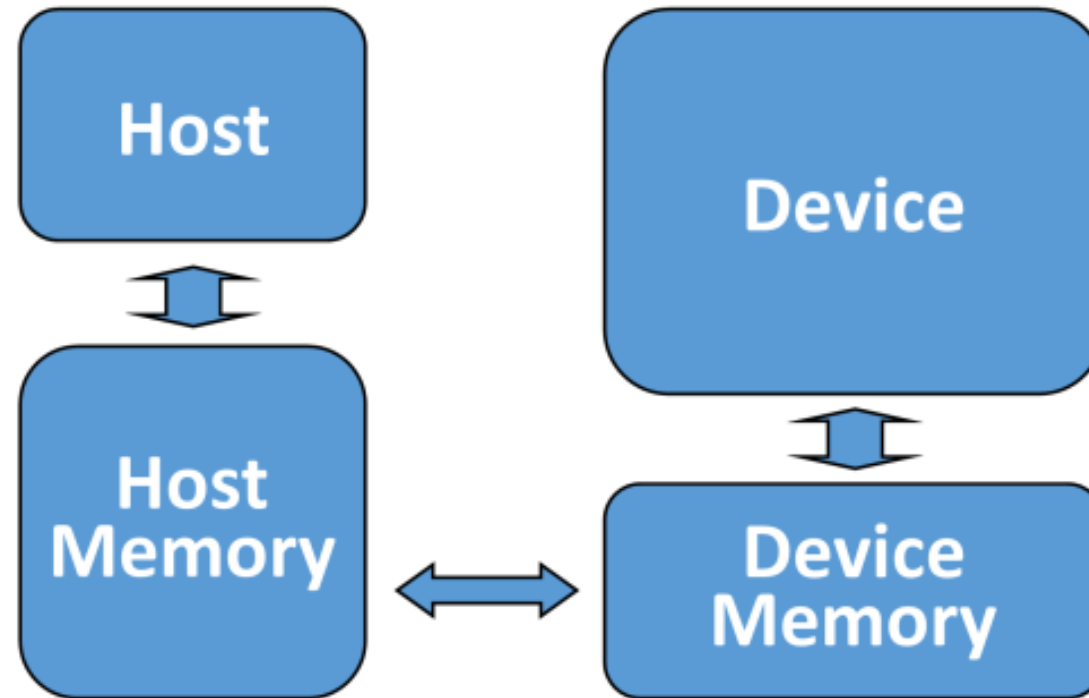
In C and C++,

- the `#pragma` directive
- provides the compiler with
- information not specified in the language.

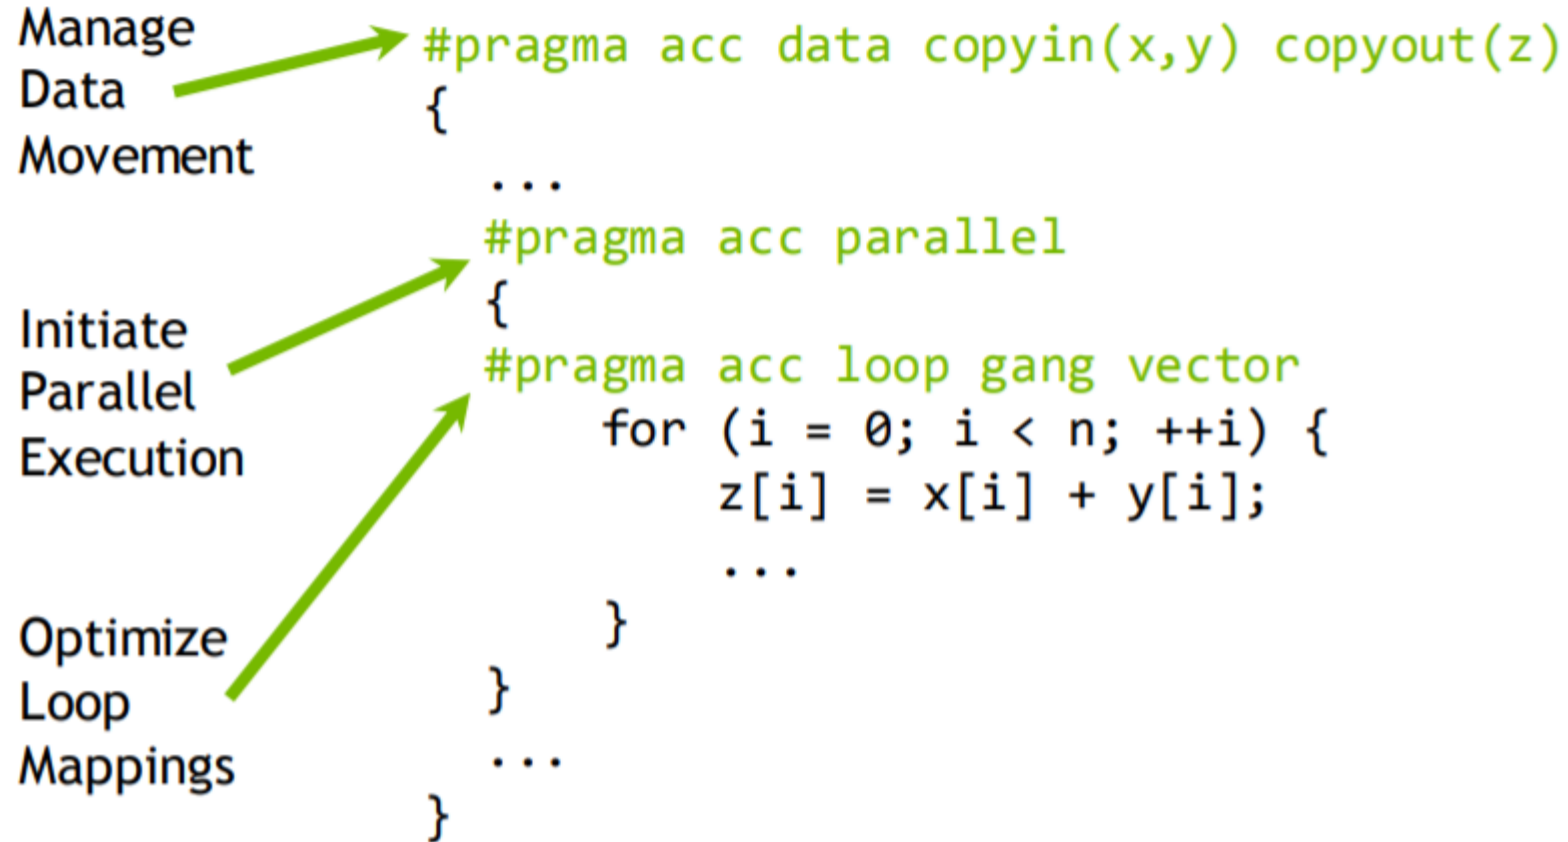
For OpenACC, they look like this:

**`#pragma acc [ the information goes here ]`**

# The OpenACC Abstract Machine Model



# The OpenACC Directives



# Simple Matrix-Matrix Multiplication in OpenACC

```
1 void computeAcc(float *P, const float *M, const float *N, int Mh, int Mw, int Nw)
2 {
3
4     #pragma acc parallel loop copyin(M[0:Mh*Mw]) copyin(N[0:Nw*Mw]) copyout(P[0:Mh*Nw])
5     for (int i=0; i<Mh; i++) {
6         #pragma acc loop
7         for (int j=0; j<Nw; j++) {
8             float sum = 0;
9             for (int k=0; k<Mw; k++) {
10                 float a = M[i*Mw+k];
11                 float b = N[k*Nw+j];
12                 sum += a*b;
13             }
14             P[i*Nw+j] = sum;
15         }
16     }
17 }
```

# Add Pragmas to Sequential Code

The **code** is

- **identical to** the **sequential** version
- **except for** the two **pragmas**
- at lines 2 and 4.

OpenACC uses the compiler directive mechanism to extend the base language.



# Simple Matrix-Matrix Multiplication in OpenACC

```
1 void computeAcc(float *P, const float *M, const float *N, int Mh, int Mw, int Nw) {
2   #pragma acc parallel loop copyin(M[0:Mh*Mw]) copyin(N[0:Nw*Mw]) copyout(P[0:Mh*Nw])
3   for (int i=0; i<Mh; i++) {
4     #pragma acc loop
5     for (int j=0; j<Nw; j++) {
6       float sum = 0;
7       for (int k=0; k<Mw; k++) {
8         float a = M[i*Mw+k];
9         float b = N[k*Nw+j];
10        sum += a*b;
11      }
12      P[i*Nw+j] = sum;
13    }
14  }
15 }
```

tells compiler

- to execute 'i' loop
- (lines 3 through 14)
- in parallel on accelerator.

copyin/copyout specify

- how matrix data
- should be transferred between memories.

# Simple Matrix-Matrix Multiplication in OpenACC

```
1 void computeAcc(float *P, const float *M, const float *N, int Mh, int Mw, int Nw) {
2   #pragma acc parallel loop copyin(M[0:Mh*Mw]) copyin(N[0:Nw*Mw]) copyout(P[0:Mh*Nw])
3   for (int i=0; i<Mh; i++) {
4     #pragma acc loop
5     for (int j=0; j<Nw; j++) {
6       float sum = 0;
7       for (int k=0; k<Mw; k++) {
8         float a = M[i*Mw+k];
9         float b = N[k*Nw+j];
10        sum += a*b;
11      }
12      P[i*Nw+j] = sum;
13    }
14  }
15 }
```

tells compiler

- to map 'j' loop
- (lines 5 through 13)
- to second level
- of parallelism on accelerator.

# Motivating Goal: One Version of Code

## OpenACC programmers

- can often start with a sequential version,
- then annotate their program with directives,
- leaving most kernel details and data transfers
- to the OpenACC compiler.

OpenACC code can be compiled by non-OpenACC compilers by ignoring the pragmas.

# Reality is More Complicated

Reality check:

- can be **difficult to write code**
- that works **correctly and well**
- **with and without pragmas.**

Some OpenACC programs

- behave differently or even incorrectly
- if pragmas are ignored.

# Pitfall: Strong Dependence on Compiler

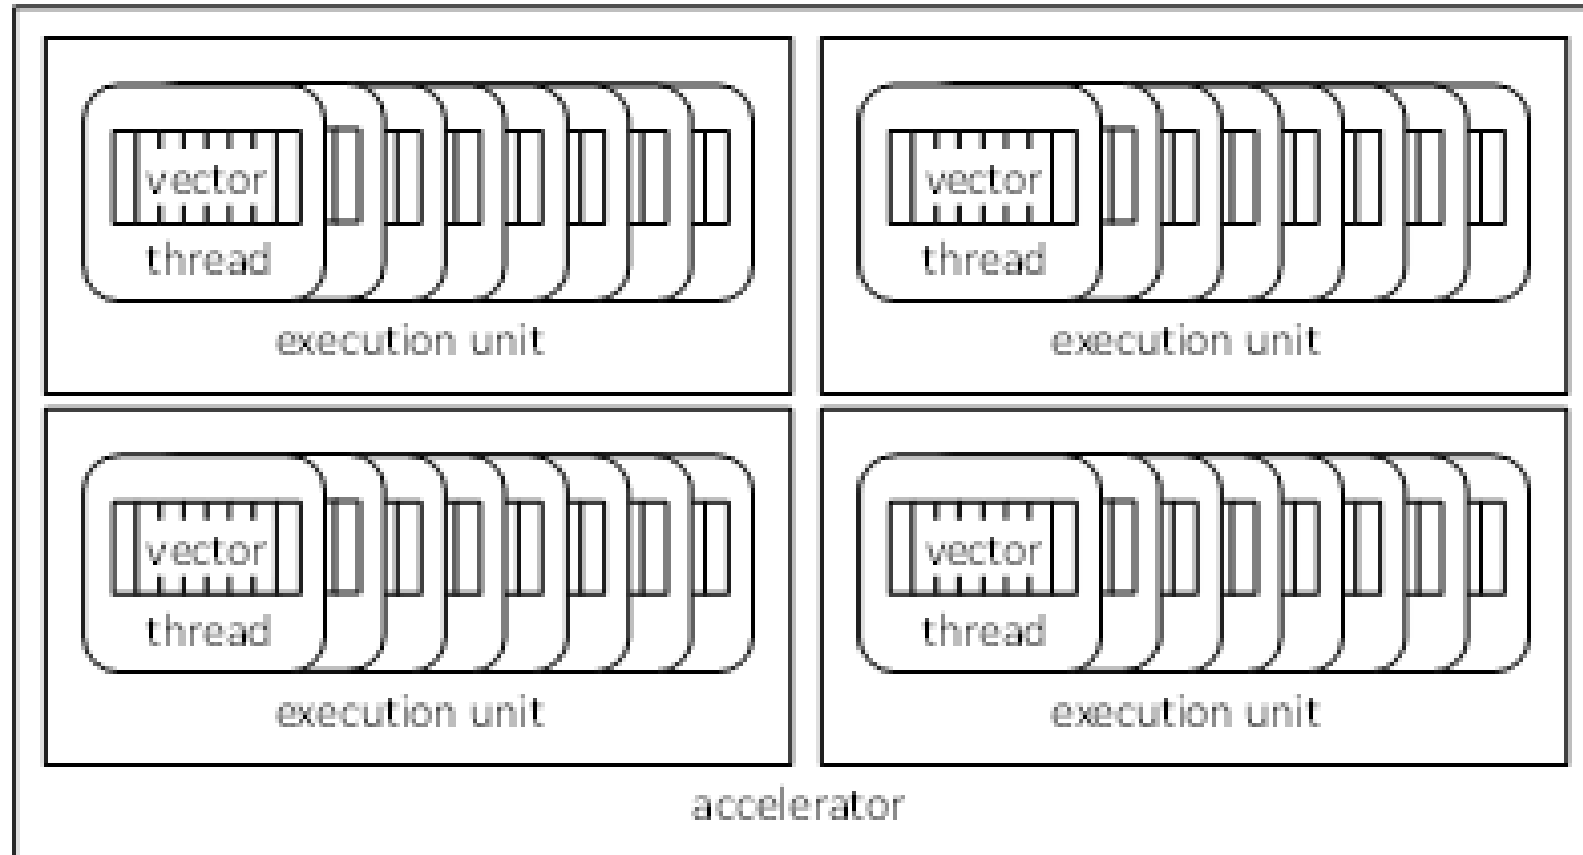
Some OpenACC pragmas

- are hints to the OpenACC compiler,
- which may or may not be able to act accordingly

**Performance depends** heavily

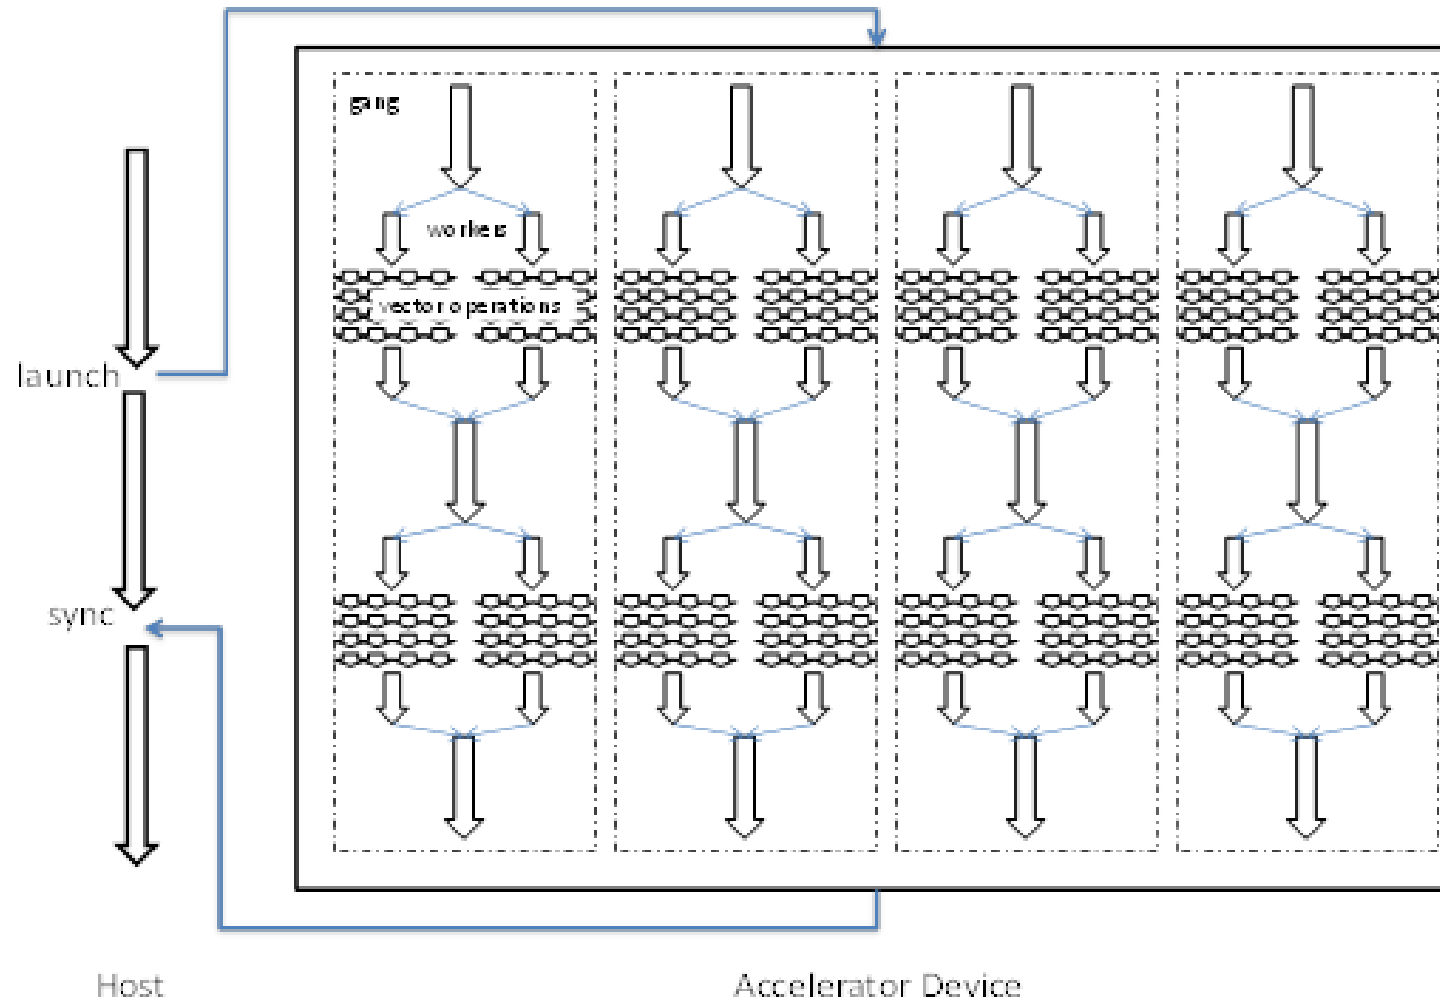
- **on the quality of the compiler**
- (more so than with CUDA or OpenCL).

# OpenACC Device Model



Currently OpenACC does not allow user-specified synchronization across threads.

# OpenACC Execution Model (Terminology: Gangs and Works)



# Parallel vs. Loop Constructs

```
#pragma acc parallel loop copyin(M[0:Mh*Mw]) copyin(N[0:Nw*Mw]) copyout(P[0:Mh*Nw])  
for (int i=0; i<Mh; i++) {  
    ...  
}
```

is equivalent to:

```
#pragma acc parallel copyin(M[0:Mh*Mw]) copyin(N[0:Nw*Mw]) copyout(P[0:Mh*Nw])  
{  
    #pragma acc loop  
    for (int i=0; i<Mh; i++) {  
        ...  
    }  
}
```

(a parallel region that consists of just a loop)



# Parallel Construct

- A parallel construct is executed on an accelerator
- One can specify the number of gangs and number of works in each gang
- Programmer's directive

```
#pragma acc parallel copyout(a) num_gangs(1024) num_workers(32)  
{  
    a = 23;  
}
```

1024\*32 workers will be created. a=23 will be executed  
redundantly by all 1024 gang leads

# What does each “Gang Loop” do?

```
#pragma acc parallel num_gangs(1024)  
{  
    for (int i=0; i<2048; i++) {  
        ...  
    }  
}
```

The for-loop will be  
redundantly executed by  
1024 gangs

```
#pragma acc parallel num_gangs(1024)  
{  
    #pragma acc loop gang  
        for (int i=0; i<2048; i++) {  
            ...  
        }  
}
```

The 2048 iterations of the  
for-loop will be divided  
among 1024 gangs for  
execution

# Worker Loop

```
#pragma acc parallel num_gangs(1024) num_workers(32)  
{  
    #pragma acc loop gang  
    for (int i=0; i<2048; i++) {  
        #pragma acc loop worker  
        for (int j=0; j<512; j++) {  
            foo(i,j);  
        }  
    }  
}
```

1024\*32=32K workers will be created, each executing  $1\text{M}/32\text{K} = 32$  instance of foo()

# A More Complex Example

```
#pragma acc parallel num_gangs(32)
```

```
{
```

```
    Statement 1; Statement 2;
```

```
    #pragma acc loop gang
```

```
    for (int i=0; i<n; i++) {
```

```
        Statement 3; Statement 4;
```

```
    }
```

```
    Statement 5; Statement 6;
```

```
    #pragma acc loop gang
```

```
    for (int i=0; i<m; i++) {
```

```
        Statement 7; Statement 8;
```

```
    }
```

```
    Statement 9;
```

```
    if (condition)
```

```
        Statement 10;
```

```
}
```

- Statements 1 and 2 are redundantly executed by 32 gangs
- The n for-loop iterations are distributed to 32 gangs

# Kernel Regions

```
#pragma acc kernels
```

```
{  
    #pragma acc loop num_gangs(1024)  
    for (int i=0; i<2048; i++) {  
        a[i] = b[i];  
    }  
    #pragma acc loop num_gangs(512)  
    for (int j=0; j<2048; j++) {  
        c[j] = a[j]*2;  
    }  
    for (int k=0; k<2048; k++) {  
        d[k] = c[k];  
    }  
}
```

- Kernel constructs are descriptive of programmer intentions (suggestions)

# Reduction

```
#pragma acc parallel loop
reduction(+:sum)
for(int i=0;i<n;i++) {
    sum +=
        xcoefs[i]*ycoefs[i];
}
```

- Because each iteration of the loop adds to the variable sum, we must declare a reduction.
- A parallel reduction may return a slightly different result than a sequential addition due to floating point limitations.

# C/C++ vs. FORTRAN

// C or C++

```
#pragma acc <directive> <clauses>  
{ ... }
```

! Fortran

```
!$acc <directive> <clauses>
```

...

```
!$acc end <directive>
```

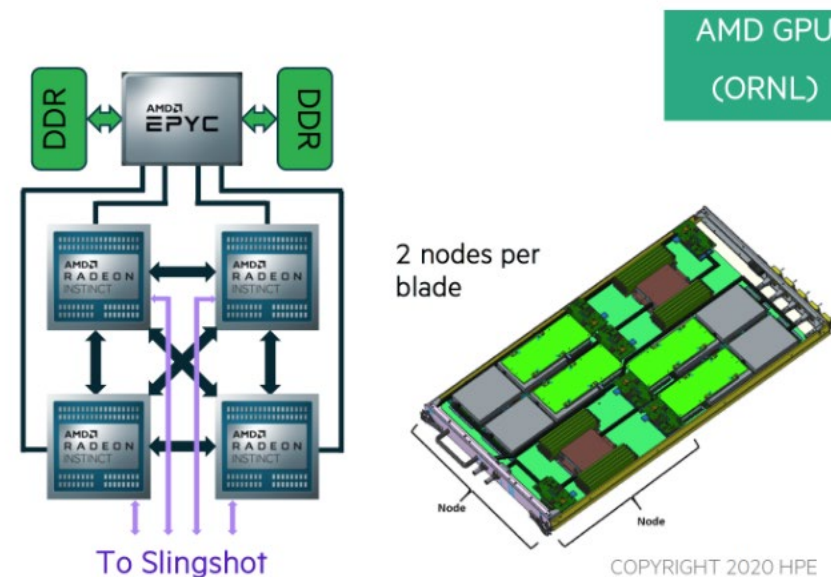
# Top 5 Supercomputers (Fall 2022)

Rank	System	Cores	Rmax (PFlop/s)	Rpeak (PFlop/s)	Power (kW)
1	<b>Frontier</b> - HPE Cray EX235a, AMD Optimized 3rd Generation EPYC 64C 2GHz, AMD Instinct MI250X, Slingshot-11, HPE DOE/SC/Oak Ridge National Laboratory United States	8,730,112	1,102.00	1,685.65	21,100
2	<b>Supercomputer Fugaku</b> - Supercomputer Fugaku, A64FX 48C 2.2GHz, Tofu interconnect D, Fujitsu RIKEN Center for Computational Science Japan	7,630,848	442.01	537.21	29,899
3	<b>LUMI</b> - HPE Cray EX235a, AMD Optimized 3rd Generation EPYC 64C 2GHz, AMD Instinct MI250X, Slingshot-11, HPE EuroHPC/CSC Finland	2,220,288	309.10	428.70	6,016
4	<b>Leonardo</b> - BullSequana XH2000, Xeon Platinum 8358 32C 2.6GHz, NVIDIA A100 SXM4 40 GB, Quad-rail NVIDIA HDR100 Infiniband, Atos EuroHPC/CINECA Italy	1,463,616	174.70	255.75	5,610
5	<b>Summit</b> - IBM Power System AC922, IBM POWER9 22C 3.07GHz, NVIDIA Volta GV100, Dual-rail Mellanox EDR Infiniband, IBM DOE/SC/Oak Ridge National Laboratory United States	2,414,592	148.60	200.79	10,096

## ORNL Frontier

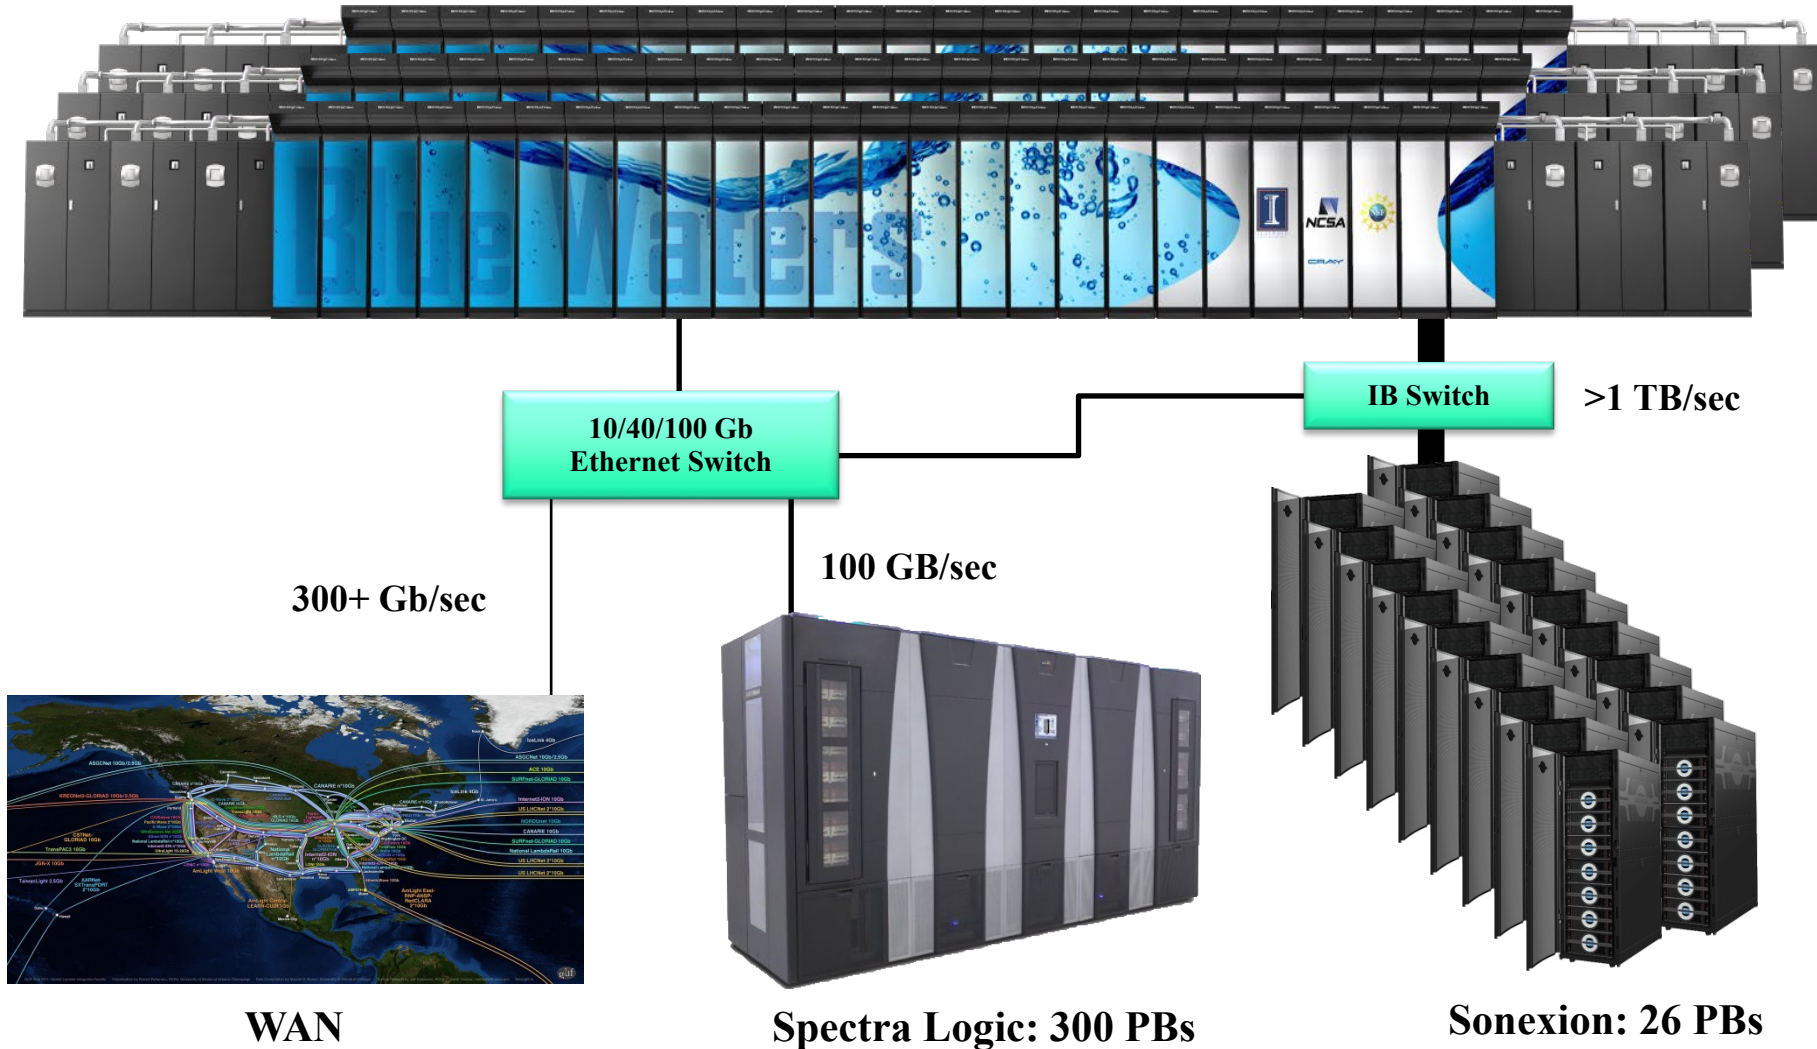
9472 AMD CPUs

37,888 AMD GPUs

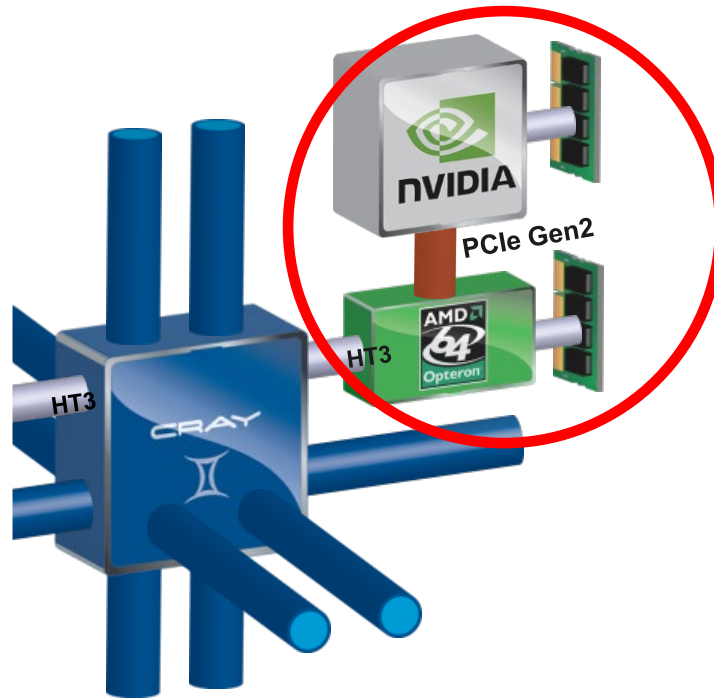




# Blue Waters @ UIUC (2013-2021)



# Cray XK7 Nodes

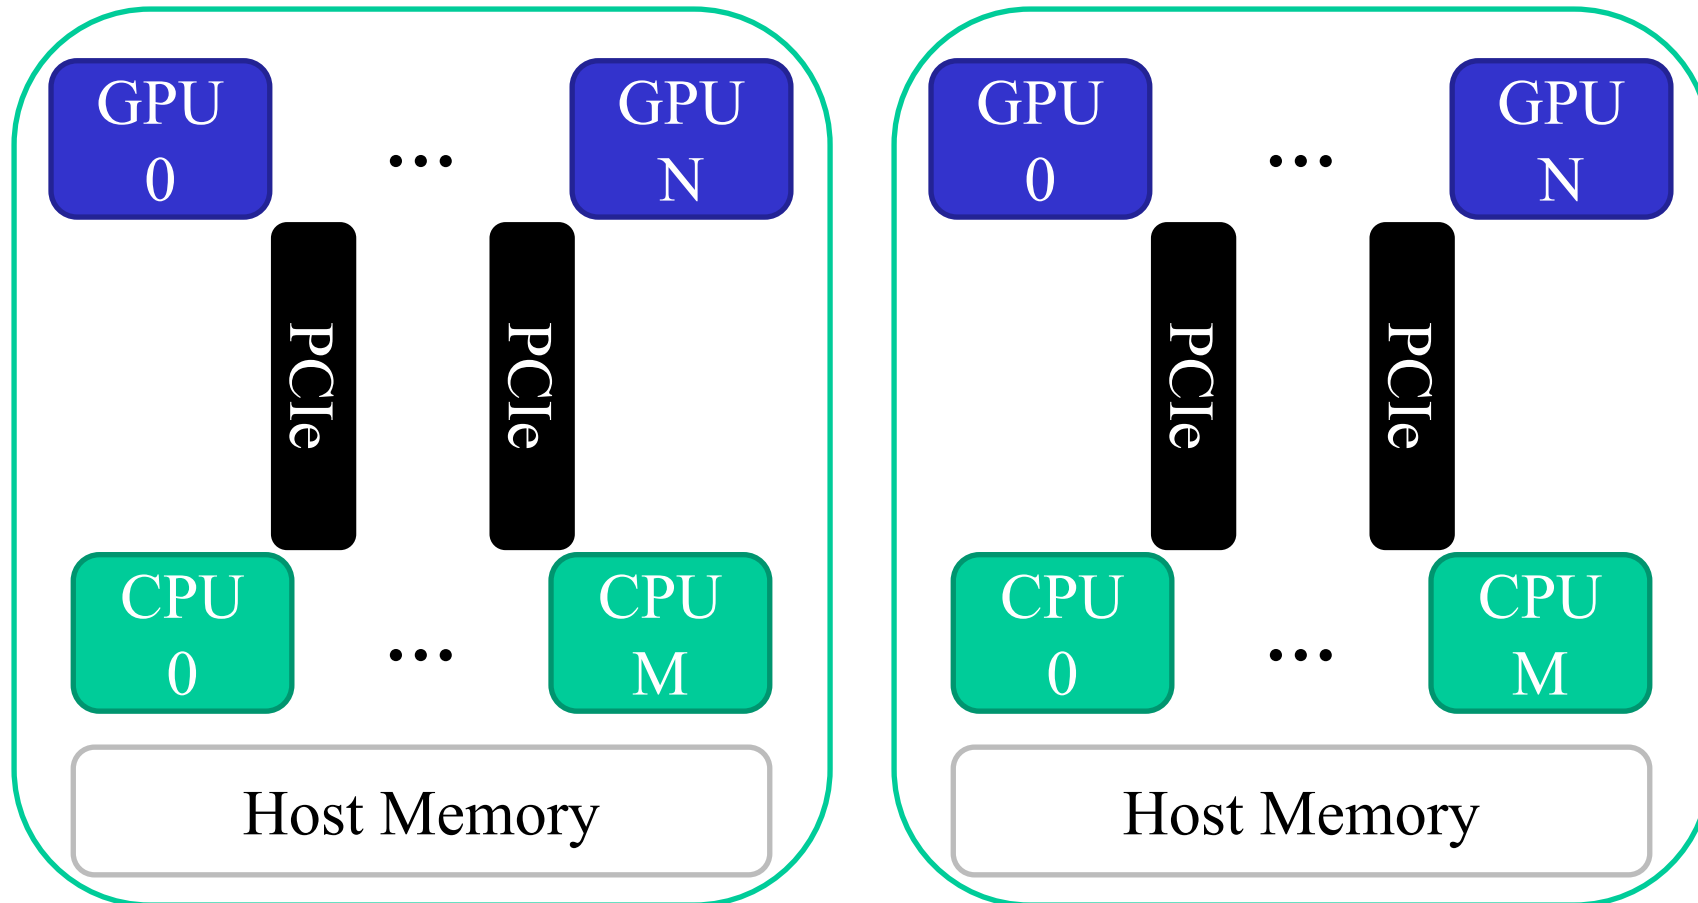


**Blue Waters contains 4,224  
Cray XK7 compute nodes.**

- Dual-socket Node
  - One AMD Interlagos chip
    - 8 core modules, 32 threads
    - 156.5 GFs peak performance
    - 32 GBs memory
      - 51 GB/s bandwidth
  - One NVIDIA Kepler chip
    - 1.3 TFs peak performance
    - 6 GBs GDDR5 memory
      - 250 GB/sec bandwidth
  - Gemini Interconnect
    - Same as XE6 nodes

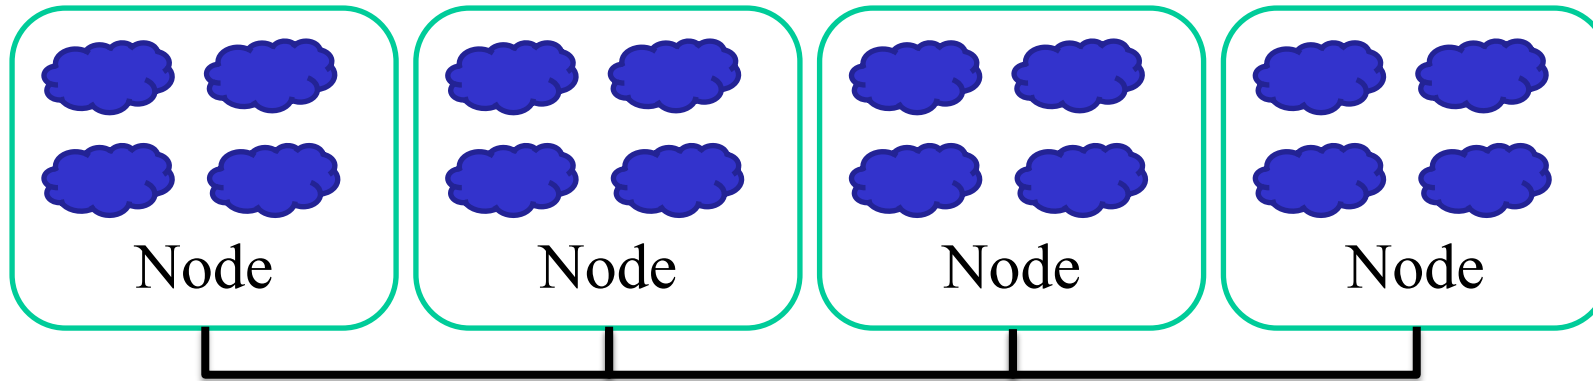
# Abstract CUDA-based Node

- Each node contains  $N$  GPUs



# MPI Model

- Many processes distributed in a cluster



- Each process computes part of the output
- Processes communicate with each other through message passing (not global memory)
- Processes can synchronize through messages

# MPI Initialization, Info

- User launches an MPI job with X processes by executing in the command shell
  - `MPRun -np X`
- `int MPI_Init(int *argc, char ***argv)`
  - Initialize MPI
- `MPI_COMM_WORLD`
  - MPI group formed with all allocated nodes
- `int MPI_Comm_rank(MPI_Comm comm, int *rank)`
  - Rank of the calling process in group of comm
- `int MPI_Comm_size(MPI_Comm comm, int *size)`
  - Number of processes in the group of comm

# Vector Addition: Main Process

```
int main(int argc, char *argv[]) {
    int vector_size = 1024 * 1024 * 1024;
    int pid=-1, np=-1;

    MPI_Init(&argc, &argv);
    MPI_Comm_rank(MPI_COMM_WORLD, &pid);
    MPI_Comm_size(MPI_COMM_WORLD, &np);

    if(np < 3) {
        if(0 == pid) printf("Needed 3 or more processes.\n");
        MPI_Abort( MPI_COMM_WORLD, 1 ); return 1;
    }
    if(pid < np - 1)
        compute_node(vector_size / (np - 1));
    else
        data_server(vector_size);

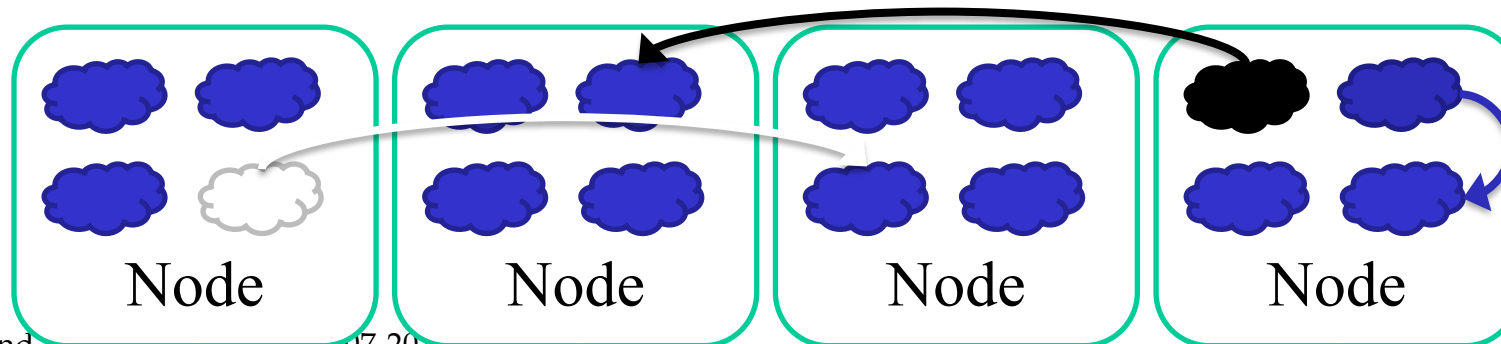
    MPI_Finalize();
    return 0;
}
```

# MPI Sending Data

- `int MPI_Send(void *buf, int count, MPI_Datatype datatype, int dest, int tag, MPI_Comm comm)`
  - `buf`: Starting address of send buffer
  - `count`: Number of elements in send buffer (nonnegative integer)
  - `datatype`: Datatype of each send buffer element
  - `dest`: Rank of destination (integer)
  - `tag`: Message tag (integer)
  - `comm`: Communicator (handle)

# MPI Sending Data

- `int MPI_Send(void *buf, int count, MPI_Datatype datatype, int dest, int tag, MPI_Comm comm)`
  - **Buf**: Initial address of send buffer
  - **Count**: Number of elements in send buffer (nonnegative integer)
  - **Datatype**: Datatype of each send buffer element
  - **Dest**: Rank of destination (integer)
  - **Tag**: Message tag (integer)
  - **Comm**: Communicator (handle)



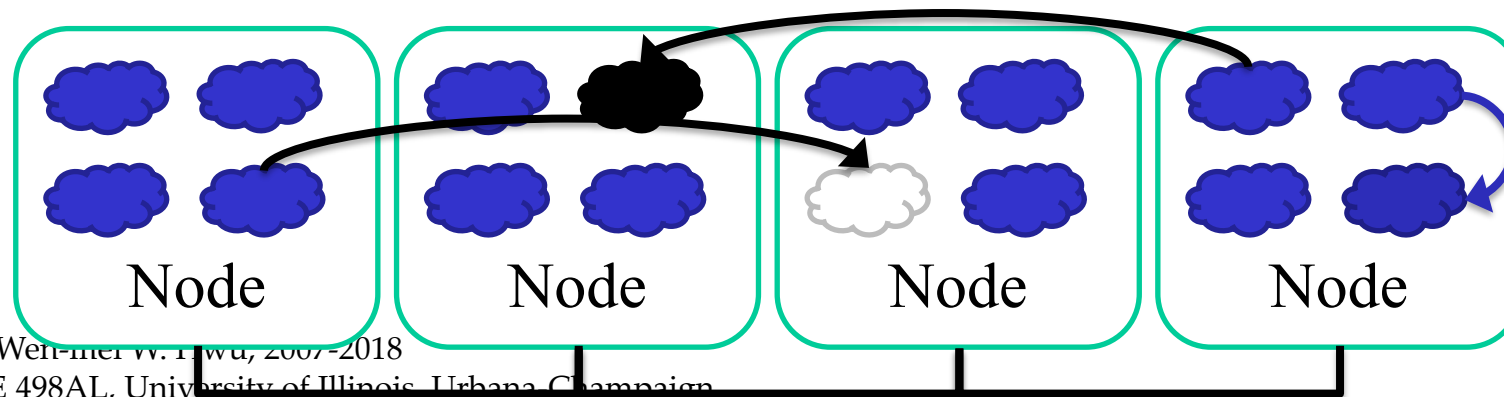


# MPI Receiving Data

- `int MPI_Recv(void *buf, int count, MPI_Datatype datatype, int source, int tag, MPI_Comm comm, MPI_Status *status)`
  - **Buf**: Starting address of receive buffer
  - **Count**: Maximum number of elements in receive buffer (non-negative integer)
  - **Datatype**: Datatype of each receive buffer element
  - **Source**: Rank of source (integer)
  - **Tag**: Message tag (integer)
  - **Comm**: Communicator (handle)
  - **Status**: Status object

# MPI Receiving Data

- `int MPI_Recv(void *buf, int count, MPI_Datatype datatype, int source, int tag, MPI_Comm comm, MPI_Status *status)`
  - **Buf**: Initial address of receive buffer
  - **Count**: Maximum number of elements in receive buffer (non-negative integer)
  - **Datatype**: Datatype of each receive buffer element
  - **Source**: Rank of source (integer)
  - **Tag**: Message tag (integer)
  - **Comm**: Communicator (handle)
  - **Status**: Status object (Status)



# Vector Addition: Server Process (I)

```
void data_server(unsigned int vector_size) {
    int np, num_nodes = np - 1, first_node = 0, last_node = np - 2;
    unsigned int num_bytes = vector_size * sizeof(float);
    float *input_a = 0, *input_b = 0, *output = 0;

    /* Set MPI Communication Size */
    MPI_Comm_size(MPI_COMM_WORLD, &np);

    /* Allocate input data */
    input_a = (float *)malloc(num_bytes);
    input_b = (float *)malloc(num_bytes);
    output = (float *)malloc(num_bytes);
    if(input_a == NULL || input_b == NULL || output == NULL) {
        printf("Server couldn't allocate memory\n");
        MPI_Abort( MPI_COMM_WORLD, 1 );
    }
    /* Initialize input data */
    random_data(input_a, vector_size , 1, 10);
    random_data(input_b, vector_size ,
```

# Vector Addition: Server Process (II)

```
/* Send data to compute nodes */
float *ptr_a = input_a;
float *ptr_b = input_b;

for(int process = 1; process < last_node; process++) {
    MPI_Send(ptr_a, vector_size / num_nodes, MPI_FLOAT,
             process, DATA_DISTRIBUTE, MPI_COMM_WORLD);
    ptr_a += vector_size / num_nodes;

    MPI_Send(ptr_b, vector_size / num_nodes, MPI_FLOAT,
             process, DATA_DISTRIBUTE, MPI_COMM_WORLD);
    ptr_b += vector_size / num_nodes;
}
```

# Vector Addition: Server Process (III)

```
/* Wait for compute to complete*/
MPI_Barrier(MPI_COMM_WORLD);

/* Collect output data */
MPI_Status status;
for(int process = 0; process < num_nodes; process++) {
    MPI_Recv(output + process * num_points / num_nodes,
             num_points / num_comp_nodes, MPI_REAL, process,
             DATA_COLLECT, MPI_COMM_WORLD, &status );
}

/* Store output data */
store_output(output, dimx, dimy, dimz);

/* Release resources */
free(input);
free(output);
}
```

# Vector Addition: Compute Process (I)

```
void compute_node(unsigned int vector_size ) {
    int np;
    unsigned int num_bytes = vector_size * sizeof(float);
    float *input_a, *input_b, *output;
    MPI_Status status;

    MPI_Comm_size(MPI_COMM_WORLD, &np);
    int server_process = np - 1;

    /* Alloc host memory */
    input_a = (float *)malloc(num_bytes);
    input_b = (float *)malloc(num_bytes);
    output = (float *)malloc(num_bytes);

    /* Get the input data from server process */
    MPI_Recv(input_a, vector_size, MPI_FLOAT, server_process,
             DATA_DISTRIBUTE, MPI_COMM_WORLD, &status);
    MPI_Recv(input_b, vector_size, MPI_FLOAT, server_process,
             DATA_DISTRIBUTE, MPI_COMM_WORLD, &status);
}
```

# Vector Addition: Compute Process (II)

```
/* Compute the partial vector addition */
for(int i = 0; i < vector_size; ++i) {
    output[i] = input_a[i] + input_b[i];
}

/* Or, can offload to GPU here */
/* cudaMalloc(), cudaMemcpy(), kernel launch, etc. */

MPI_Barrier(MPI_COMM_WORLD);

/* Send the output */
MPI_Send(output, vector_size, MPI_FLOAT,
         server_process, DATA_COLLECT, MPI_COMM_WORLD);

/* Release memory */
free(input_a);
free(input_b);
free(output);
}
```



**ANY MORE QUESTIONS?**  
**READ CHAPTER 15**

**Also see <https://developer.nvidia.com/intro-to-openacc-course-2016>**